

# Lab 6 (Bayesian Estimation and MCMC)

Shreya Rao sr3843

06/10/2021

## Goals

This lab has two goals. The first goal is to use the **Accept-Reject** algorithm to simulate from a mixture of two normals. The second goal is to utilize Bayesian methods and the famous **Markov Chain Monte Carlo** algorithm to estimate the mixture parameter  $\delta$ .

## Background: (Mixture)

A mixture distribution is the probability distribution of a random variable that is derived from a collection of other random variables (Wiki). In our case we consider a mixture of two normal distributions. Here we assume that our random variable is governed by the probability density  $f(x)$ , defined by

$$\begin{aligned} f(x) &= f(x; \mu_1, \sigma_1, \mu_2, \sigma_2, \delta) \\ &= \delta f_1(x; \mu_1, \sigma_1) + (1 - \delta) f_2(x; \mu_2, \sigma_2) \\ &= \delta \frac{1}{\sqrt{2\pi\sigma_1^2}} \exp -\frac{1}{2\sigma_1^2}(x - \mu_1)^2 + (1 - \delta) \frac{1}{\sqrt{2\pi\sigma_2^2}} \exp -\frac{1}{2\sigma_2^2}(x - \mu_2)^2, \end{aligned}$$

where  $-\infty < x < \infty$  and the parameter space is defined by  $-\infty < \mu_1, \mu_2 < \infty$ ,  $\sigma_1, \sigma_2 > 0$ , and  $0 \leq \delta \leq 1$ . The **mixture parameter**  $\delta$  governs how much mass gets placed on the first distribution  $f(x; \mu_1, \sigma_1)$  and the complement of  $\delta$  governs how much mass gets placed on the other distribution  $f_2(x; \mu_2, \sigma_2)$ .

To further motivate this setting, consider simulating  $n = 10,000$  heights from the population of both males and females. Assume that males are distributed normal with mean  $\mu_1 = 70$ [in] and standard deviation  $\sigma_1 = 3$ [in] and females are distributed normal with mean  $\mu_2 = 64$ [in] and standard deviation  $\sigma_2 = 2.5$ [in]. Also assume that each distribution contributes equal mass, i.e., set the mixture parameter to  $\delta = .5$ . The distribution of males is governed by

$$f_1(x; \mu_1, \sigma_1) = \frac{1}{\sqrt{2\pi\sigma_1^2}} \exp -\frac{1}{2\sigma_1^2}(x - \mu_1)^2, \quad -\infty < x < \infty,$$

and the distribution of females is governed by

$$f_2(x; \mu_2, \sigma_2) = \frac{1}{\sqrt{2\pi\sigma_2^2}} \exp -\frac{1}{2\sigma_2^2}(x - \mu_2)^2, \quad -\infty < x < \infty.$$


Below shows the pdf of  $f_1(x; \mu_1, \sigma_1)$ ,  $f_2(x; \mu_2, \sigma_2)$  and the mixture  $f(x)$  all on the same plot.

```
x <- seq(45, 90, by = 0.05)
n.x <- length(x)
f_1 <- dnorm(x, mean = 70, sd = 3)
```

```
f_2 <- dnorm(x, mean = 64, sd = 2.5)
f <- function(x) {
  return(0.5 * dnorm(x, mean = 70, sd = 3) + 0.5 * dnorm(x,
    mean = 64, sd = 2.5))
}

plot_df <- data.frame(x = c(x, x, x), Density = c(f_1, f_2, f(x)),
  Distribution = c(rep("Male", n.x), rep("Female", n.x), rep("Mixture",
    n.x)))

library(ggplot2)
ggplot(data = plot_df) + geom_line(mapping = aes(x = x, y = Density,
  color = Distribution)) + labs(title = "Mixture of Normals")
```



## Part I: Simulating a Mixture of Normals

The first goal is to simulate from the mixture distribution

$$\delta f_1(x; \mu_1, \sigma_1) + (1 - \delta) f_2(x; \mu_2, \sigma_2),$$

where  $\mu_1 = 70, \sigma_1 = 3, \mu_2 = 64, \sigma_2 = 2.5, \delta = .5$ . We use the accept-reject algorithm to accomplish this task.

First we must choose the “easy to simulate” distribution  $g(x)$ . For this problem choose  $g(x)$  to be a Cauchy distribution centered at 66 with scale parameter 7.

```
g <- function(x) {
  s = 7
  l = 66
  return(1/(pi * s * (1 + ((x - l)/s)^2)))
}
```

### Perform the following tasks

- 1) Identify a **suitable** value of  $\alpha$  such that your envelope function  $e(x)$  satisfies


$$f(x) \leq e(x) = g(x)/\alpha, \text{ where } 0 < \alpha < 1.$$

Note that you must choose  $\alpha$  so that  $e(x)$  is close to  $f(x)$ . There is not one unique solution to this problem. The below plot shows how  $\alpha = .20$  creates an envelope function that is too large. Validate your choice of  $\alpha$  with a graphic similar to below.

```
# Choose alpha
alpha <- 0.2

# Define envelope e(x)
e <- function(x) {
  return(g(x)/alpha)
}

# Plot
x.vec <- seq(30, 100, by = 0.1)
ggplot() + geom_line(mapping = aes(x = x.vec, y = f(x.vec)),
  col = "purple") + geom_line(mapping = aes(x = x.vec, y = e(x.vec)),
  col = "green")
```




```
# Is  $g(x) > f(x)$ ?
all(e(x.vec) > f(x.vec))
```

```
## [1] TRUE
```

### Solution

```
alpha <- 0.42

# Plot
x.vec <- seq(30, 100, by = 0.1)
ggplot() + geom_line(mapping = aes(x = x.vec, y = f(x.vec)),
  col = "purple") + geom_line(mapping = aes(x = x.vec, y = e(x.vec)),
  col = "green")
```



```
# Is  $g(x) > f(x)$ ?
all(e(x.vec) > f(x.vec))
```

```
## [1] TRUE
```

- 2) Write a function named **r.norm.mix()** that simulates **n.samps** from the normal-mixture  $f(x)$ . To accomplish this task you will wrap a function around the accept-reject algorithm from the lecture notes. Also include the acceptance rate, i.e., how many times did the algorithm accept a draw compared to the total number of trials performed. Your function should return a list of two elements: (i) the simulated vector mixture and (ii) the proportion of accepted cases. Run your function **r.norm.mix()** to simulate 10,000 cases and display the first 20 values. What's the proportion of accepted cases? Compare this number to your chosen  $\alpha$  and comment on the result. The code below should help you get started.

### Solution

```
r.norm.mix <- function(n.samps) {
  n <- 0 # counter for number samples accepted
  m <- 0 # counter for number of trials
  samps <- numeric(n.samps) # initialize the vector of output
  while (n < n.samps) {
    m <- m + 1
    y <- rcauchy(1, location = 66, scale = 7) #sample from g
    u <- runif(1)
    if (u < f(y)/e(y)) {
      n <- n + 1
    }
  }
  list(samps[1:n], n/m)
```

```

      samps[n] <- y
    }
  }
  return(list(x = samps, alpha.hat = n.samps/m))
}

```

```

sim.norm.mix <- r.norm.mix(n.samps = 10000)
head(sim.norm.mix$x, 20)

```

```

## [1] 61.80217 67.84918 71.63273 67.95398 66.68302 74.59870 60.88830 66.11376
## [9] 64.63964 68.63727 69.15038 62.64270 63.19513 66.18190 66.31556 61.63554
## [17] 71.46076 66.73278 72.05453 66.18352

```

```
sim.norm.mix$alpha.hat
```

```
## [1] 0.4193575
```

```
alpha
```

```
## [1] 0.42
```

Since  $\alpha = 0.42$ , we need to expect to accept around 42% of the cases, which is corroborated by `alpha.hat`. As  $\alpha$  gets larger,  $e(x)$  moves closer to  $f(x)$  and % of cases in `alpha.hat` increases.

- 3) Using **ggplot** or **base R**, construct a histogram of the simulated mixture distribution with the true mixture pdf  $f(x)$  overlayed on the plot.


## Solution

```


ggplot() + geom_histogram(aes(x = sim.norm.mix$x, y = ..density..)) +
  geom_line(aes(x = x, y = f(x)), col = "pink")

```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



```
hist(sim.norm.mix$x, probability = T)
lines(x, f(x), col = "purple")
```



## Part II: Bayesian Statistics and MCMC

Suppose that the experimenter collected 100 cases from the true mixture-normal distribution  $f(x)$ . To solve problems (4) through (8) we analyze one realized sample from our function `r.norm.mix()`. In practice this dataset would be collected and not simulated. Uncomment the below code to simulate our dataset `x`. If you failed to solve Part I, then read in the csv file `mixture_data.csv` posted on Canvas.

### Solution


```
# Simulate data set.seed(1983) x <- r.norm.mix(n.samps=100)$x
# head(x) hist(x,breaks=20,xlab='X',main='')

# Or read data
x <- read.csv("mixture_data.csv")$x
head(x)
```

```
## [1] 71.66666 63.91096 67.06554 65.49516 70.34363 65.69982
```

```
hist(x, breaks = 20, xlab = "X", main = "")
```





Further, suppose that we know the true heights and standard deviations of the two normal distributions but the mixture parameter  $\delta$  is unknown. In this case, we know  $\mu_1 = 70$ ,  $\sigma_1 = 3$ ,  $\mu_2 = 64$ ,  $\sigma_2 = 2.5$ . The goal of this exercise is to utilize **maximum likelihood** and **MCMC Bayesian techniques** to estimate mixture parameter  $\delta$ .

### Maximum likelihood Estimator of Mixture Parameter

- 4) Set up the likelihood function  $L(\delta|x_1, \dots, x_{100})$  and define it as **mix.like()**. The function should have two inputs including the parameter **delta** and data vector **x**. Evaluate the likelihood at the parameter values **delta=.2**, **delta=.4**, and **delta=.6**. Note that all three evaluations will be very small numbers. Which delta ( $\delta = .2, .4, .6$ ) is the most likely to have generated the dataset **x**?

**Solution**

```
mix.like <- function(delta, x) {
  ll.func <- prod(delta * dnorm(x, mean = 70, sd = 3) + (1 -
    delta) * dnorm(x, mean = 64, sd = 2.5))
  return(ll.func)
}

mix.like(0.2, x = x)
```

```
## [1] 4.546128e-129
```

```
mix.like(0.4, x = x)
```

```
## [1] 1.967793e-124
```

```
mix.like(0.6, x = x)
```

```
## [1] 2.838963e-125
```


delta = 0.4 is most likely

- 5) Compute the maximum likelihood estimator of mixture parameter  $\delta$ . To accomplish this task, apply your likelihood function **mix.like()** across the vector **seq(.1,.99,by=.001)**. The solution to this exercise is given below.

```
# apply likelihood function across seq vector 0.1-0.99
delta <- seq(0.1, 0.99, by = 0.001)
MLE.values <- sapply(delta, mix.like, x = x)


# which delta value has the max likelihood?
delta.MLE <- delta[which.max(MLE.values)]

# plotting MLE values with max MLE value
plot(delta, MLE.values, ylab = "Likelihood", type = "l")
abline(v = delta.MLE, col = "blue")
text(x = delta.MLE + 0.08, y = mix.like(delta = 0.45, x = x),
     paste(delta.MLE), col = "blue")
```



Optional:

```
neg.mix.ll <- function(delta, data = x) {  
  return(-1 * sum(log(delta * dnorm(x, mean = 70, sd = 3) +  
    (1 - delta) * dnorm(x, mean = 64, sd = 2.5))))  
}  
  
plot(delta, sapply(delta, neg.mix.ll, data = x), type = "l")
```



```
nlm(neg.mix.ll, p = 0.5)
```

```
## Warning in log(delta * dnorm(x, mean = 70, sd = 3) + (1 - delta) * dnorm(x, :  
## NaNs produced
```

```
## Warning in nlm(neg.mix.ll, p = 0.5): NA/Inf replaced by maximum positive value
```

```
## Warning in log(delta * dnorm(x, mean = 70, sd = 3) + (1 - delta) * dnorm(x, :  
## NaNs produced
```

```
## Warning in nlm(neg.mix.ll, p = 0.5): NA/Inf replaced by maximum positive value
```

```
## $minimum  
## [1] 284.3022  
##
```

```
## $estimate
## [1] 0.4632656
##
## $gradient
## [1] 0.0001152785
##
## $code
## [1] 1
##
## $iterations
## [1] 3
```

## MCMC

6) Run the Metropolis-Hastings algorithm to estimate mixture parameter  $\delta$ . In this exercise you will assume a  $\text{Beta}(\alpha = 10, \beta = 10)$  prior distribution on mixture parameter  $\delta$ . Some notes follow:

- Run 20000 iterations. I.e., simulate 20000 draws of  $\delta^{(t)}$
- Proposal distribution  $\text{Beta}(\alpha = 10, \beta = 10)$
- Independence chain with Metropolis-Hastings ratio:

$$R(\delta^{(t)}, \delta^*) = \frac{L(\delta^* | x_1, \dots, x_{100})}{L(\delta^{(t)} | x_1, \dots, x_{100})}$$


Display the first 20 simulated cases of  $\delta^{(t)}$ .

## Solution

```
# draw initial theta (delta_0 in this case)
delta_0 <- rbeta(1, shape1 = 10, shape2 = 10)
n.samps <- 20000 #number of iterations
delta_vec <- rep(NA, n.samps + 1)
delta_vec <- delta_0

# MCMC loop
for (t in 1:n.samps) {
  delta_star <- rbeta(1, shape1 = 10, shape2 = 10) #draw delta* from proposal
  delta_t <- delta_vec[t] #delta_(t)
  # compute MH ratio
  MH_ratio <- mix.like(delta = delta_star, x = x)/mix.like(delta = delta_t,
    x = x)
  # select new case
  prob_vec <- c(min(MH_ratio, 1), 1 - min(MH_ratio, 1))
  delta_vec[t + 1] <- sample(c(delta_star, delta_t), 1, prob = prob_vec)
}


hist(delta_vec, breaks = 30)
```



- 7) Construct a lineplot of the simulated Markov chain from exercise (6). The vertical axis is the simulated chain  $\delta^{(t)}$  and the horizontal axis is the number of iterations.

#### Solution

```
# lineplot of simulated Markov chain (vertical axis is the
# simulated chain)
plot(delta_vec, type = "l", xlab = "t", ylab = expression(delta[t]))
```




- 8) Plot the empirical autocorrelation function of your simulated chain  $\delta^{(t)}$ . I.e., run the function `acf()`. A quick decay of the chain's autocorrelations indicate good mixing properties.

#### Solution

```
# Plotting the empirical autocorrelation function of  
# simulated chain  
acf(delta_vec, main = "ACF: Prior Beta(10,10)")
```

### ACF: Prior Beta(10,10)



- 9) Compute the empirical Bayes estimate  $\hat{\delta}_B$  of the simulated posterior distribution  $\pi(\delta|x_1, \dots, x_n)$ . To solve this problem, simply compute the sample mean of your simulated chain  $\delta^{(t)}$  after discarding a 20% burn-in.

#### Solution

```
# discard the first 20% of the cases subset delta_vec from
# 4001 to 20000
delta_vec_new <- delta_vec[(floor(0.2 * length(delta_vec)) +
  1):length(delta_vec)]

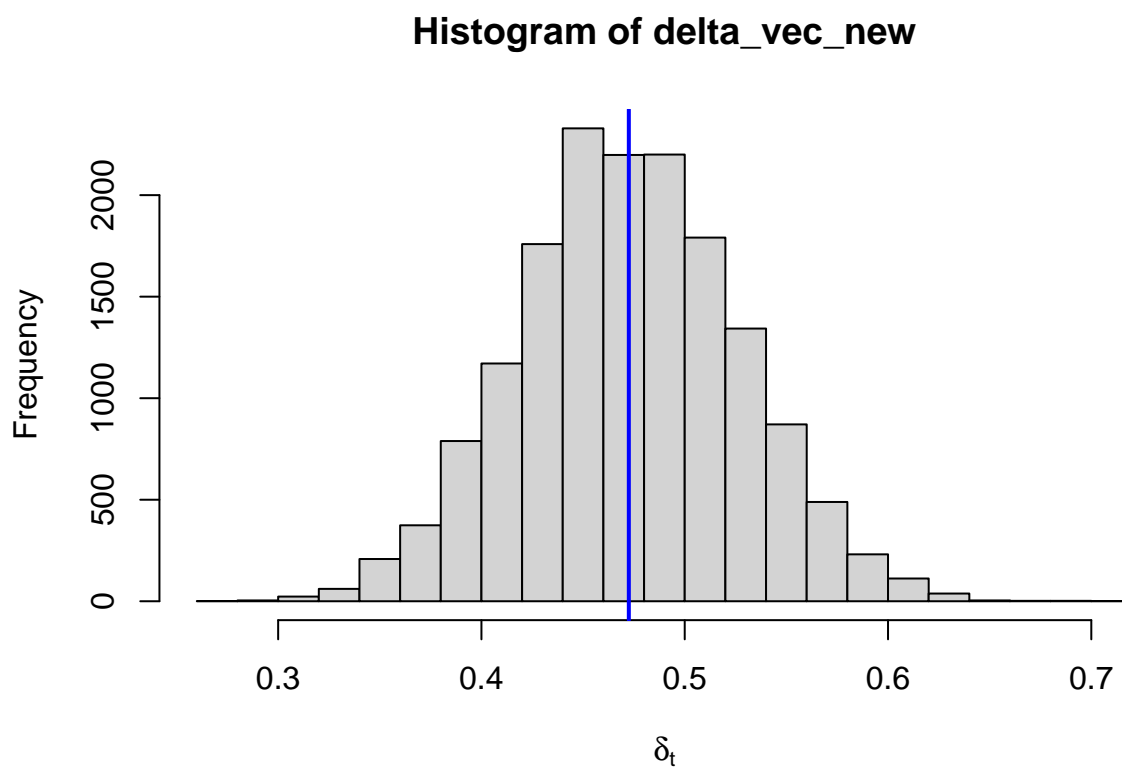
mean(delta_vec_new)
```

```
## [1] 0.4725019
```

- 10) Construct a histogram of the simulated posterior  $\pi(\delta|x_1, \dots, x_n)$  after discarding a 20% burn-in.

#### Solution

```
# histogram of the simulated posterior after discarding a 20%
# burn-in
hist(delta_vec_new, breaks = 30, xlab = expression(delta[t]))
abline(v = mean(delta_vec_new), col = "blue", lwd = 2)
```

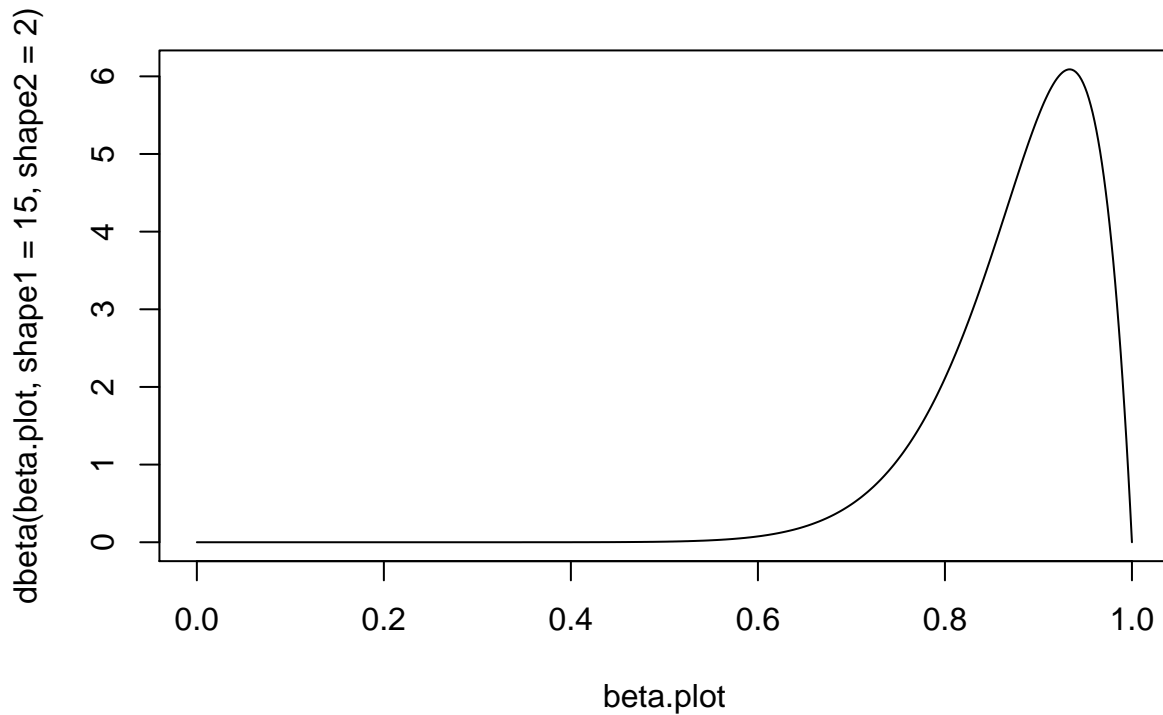


- 11) Run the Metropolis-Hastings algorithm to estimate the mixture parameter  $\delta$  using a  $\text{Beta}(\alpha = 15, \beta = 2)$  prior distribution on mixture parameter  $\delta$ . Repeat exercises 6 through 10 using the updated prior.

#### Solution

```
beta.plot <- seq(0, 1, length = 500)
plot(beta.plot, dbeta(beta.plot, shape1 = 15, shape2 = 2), type = "l")
```






```
delta_0 <- rbeta(1, shape1 = 15, shape2 = 2)
n.samps <- 20000
delta_vec <- rep(NA, n.samps + 1)
delta_vec = delta_0

for (t in 1:n.samps) {
  delta_star <- rbeta(1, shape1 = 15, shape2 = 2)
  delta_t <- delta_vec[t]
  MH_ratio <- mix.like(delta = delta_star, x = x)/mix.like(delta = delta_t,
    x = x)
  prob_vec <- c(min(1, MH_ratio), 1 - min(1, MH_ratio))
  delta_vec[t + 1] <- sample(c(delta_star, delta_t), 1, prob = prob_vec)
}
```

### lineplot:

Construct a lineplot of the simulated Markov chain from exercise (6). The vertical axis is the simulated chain  $\delta^{(t)}$  and the horizontal axis is the number of iterations.

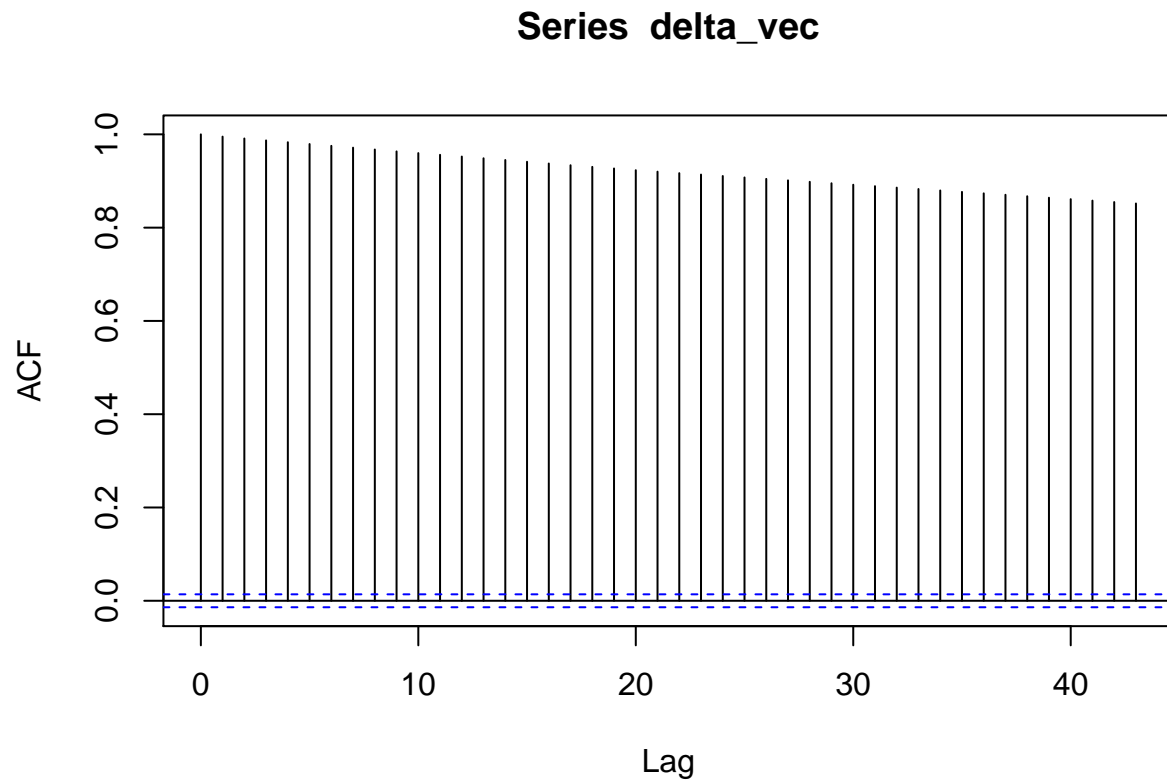
```
plot(delta_vec, type = "l", xlab = "t", ylab = expression(delta[t]))
```



#### ACF:

Plot the empirical autocorrelation function of your simulated chain  $\delta^{(t)}$ . I.e., run the function `acf()`. A slow decay of the chain's autocorrelations indicate poor mixing properties.

```
acf(delta_vec)
```



#### Bayes estimate:

Compute the empirical Bayes estimate  $\hat{\delta}_B$  of the simulated posterior distribution  $\pi(\delta|x_1, \dots, x_n)$ . To solve this problem, simply compute the sample mean of your simulated chain  $\delta^{(t)}$  after discarding a 20% burn-in. Your answer should be close to the MLE.

#### Solution

```
delta_vec_80 <- delta_vec[(floor(0.2 * length(delta_vec)) + 1):length(delta_vec)]
mean(delta_vec_80)
```

```
## [1] 0.5535371
```

MLE value is 0.463

**Posterior:** Construct a histogram of the simulated posterior  $\pi(\delta|x_1, \dots, x_n)$  after discarding a 20% burn-in.

#### Solution

```
hist(delta_vec_80)
```

**Histogram of delta\_vec\_80**

